

CSci 4061

Introduction to Operating Systems

File Systems: Basics

Chapter 5

File as Abstraction

- Container for related information
- Named
- Associated attributes

Persistence

Process abstraction

Container for a program

I/O abstraction

Source/ sink for data (using fd's FILE*)


Naming a File

```
creat/open ("path/name", ...);
```


Links: files with multiple names

Each name is an alias

```
#include <unistd.h>
```

```
 int link (const char *original_path,  
          const char *new_path)
```

cannot exist as
a file already



```
link ("foo", "bar"); // "bar" refers to file "foo"
```

```
 unlink ("bar"); // remove name "bar"
```

```
// if file is open by someone, will not actually get deleted until  
// all fd's to it are closed
```


What if unlink 'foo' here?

A; sort of like removing last pointer to object
(indirect deletion)

File Attributes: Access to metadata

 `#include <sys/stat.h>`

`int fstat (int filedes, struct stat *buf)`

 `int stat (const char *pathname,
struct stat *buf)`

MUST open file *

← stat need to open file

Structure contains file/directory info:

```
off_t st_size;      // file size
nlink_t st_nlink;   // links
mode_t st_mode;     // type + permission
time_t st_mtime;    // last modification time
```

 `fcntl` can also be used to set or get lower-level attrs

Exercise: Metadata

sleep(60)

- Write a program that monitors a given file every minute and if the file size has changed, it outputs the new size to stdout

```
#include <sys/stat.h>
int stat (const char *pathname,
struct stat *buf)
```

Structure contains file/directory level info:

```
off_t st_size;      // file size
nlink_t st_nlink;   // links
...
```

SLIDE INSERT

(backup code here)

```


:
struct stat sb;
off_t psize = -99;
while (1) {
    stat(fname, &sb);
    if (sb.st_size != psize) {
        printf("File size is now %d\n", sb.st_size);
        psize = sb.st_size;
    }
}
}

```

Storing File Meta-data: Unix inode

Allows fast for common case which is smaller file sizes

While also flexible for the less common case which is larger file sizes



`ls -li <file>` shows inode #

Filesystem

- Directory is a file as well
 - it has an **inode**
 - what are file contents?
 - list of file_name, inode pairs
- Filesystem
 - Files
 - Directories
 - Free disk sectors (free list)
 - Root dir

Filesystem (cont'd)

- On-disk organization

- `inode` for root dir of filesystem `/` stored in well-known sector on the disk
- `inode` for disk sector *free-list* also stored in a well-known sector on the disk
- Inode table or file (inode #, sector)
- These are stored in the *superblock*



Unix file types/modes

- Indicated by the first character in `ls -l`
 - `-` regular file
 - `d` directory
 - `c` character special file
 - `b` block special file
 - `p` pipe
 - `s` socket
 - `l` symbolic link



File types

- Within `stat` structure:

```
struct_t stat st;  
stat ("foo", &st);
```

Macros:

```
int S_ISDIR (st.st_mode);  
int S_ISREG (st.st_mode);  
int S_ISSOCK (st.st_mode);  
...
```

P. 158

Another look at `ls -l`

Example:


all from inode

```
drwx-xr-x  3  jon  fac 4066 Nov 2 09:14 st
```

file type

hard
links

allocation
size



Filesystem semantics: Unix

- Two processes open the same file
- Reader sees most recent write
- One reader and one writer - run together
 - File “foo” contains
“aaaaaaaaaaaaaaaaaaaaaa”

Filesystem semantics (cont'd)

```
// reader.c
#include <stdio.h>
#include <unistd.h>
#include <fcntl.h>
```

```
void main () {
    int fd, n;

    char c, buf[100];
    read (0, &c, 1);
    fd = open ("foo", O_RDONLY);
    n = read (fd, buf, 10);
    buf[n] = '\0';
    printf ("buf=%s\n", buf);
    ...
}
```

```
// writer.c
#include <stdio.h>
#include <unistd.h>
#include <fcntl.h>
```

```
void main () {
    int fd, n;

    char c, buf[100] = "bbbbbbbbbbb...";
    read (0, &c, 1);
    fd = open ("foo", O_WRONLY);
    write (fd, buf, 10);
    read (0, &c, 1);
    close (fd) }
```

Missing



Operations (r, w, x): read, write, execute

Subjects (**u**: user/owner, **g**: group, **o**: others)

Users may belong to any number of groups
(type `groups` at the shell)

(type groups at the shell)

↑
terminal

↑
terminal

File permissions (cont'd)

```
shell> ls -l
```

drwxr-xr-w 3 jon fac 46
 u g o owner group

owner group

user-group	other
1	1
2	1
3	1
4	1
5	1
6	1
7	1
8	1
9	1
10	1
11	1
12	1
13	1
14	1
15	1
16	1
17	1
18	1
19	1
20	1
21	1
22	1
23	1
24	1
25	1
26	1
27	1
28	1
29	1
30	1
31	1
32	1
33	1
34	1
35	1
36	1
37	1
38	1
39	1
40	1
41	1
42	1
43	1
44	1
45	1
46	1
47	1
48	1
49	1
50	1
51	1
52	1
53	1
54	1
55	1
56	1
57	1
58	1
59	1
60	1
61	1
62	1
63	1
64	1
65	1
66	1
67	1
68	1
69	1
70	1
71	1
72	1
73	1
74	1
75	1
76	1
77	1
78	1
79	1
80	1
81	1
82	1
83	1
84	1
85	1
86	1
87	1
88	1
89	1
90	1
91	1
92	1
93	1
94	1
95	1
96	1
97	1
98	1
99	1
100	1

When a file is created it is given a restricted permission and a default group

You can broaden or further restrict permissions

File permissions (cont'd)

```
#include <sys/types.h>
```

```
#include <sys/stat.h>
```

```
int chmod (char *path, mode_t mode);
```

prefer symbolic flags: man fstat, e.g. S_IRGRP : GRouP has Read)

0077

each octal digit is rwx (or 1) for ugo

000, 111, 111

Also at the command-line (absolute)

```
chmod 0077 st.txt
```

Also at the command-line (relative)

```
chmod go-xr st.txt
```

```
chmod u+xrw st.txt
```

Power of IDs

- Real user-id: user that actually initiated a process
 - Not executable owner!
- Effective user-id: user that system associates with the process for purposes of protection
 - Usually the same as the real user-id: this would be?
 - Sometimes want effective user-id to that of the file owner and not the user ... why?

Power of IDs (cont'd)

- How do to it?

```
shell> chmod u+s my_file
```

```
-r-s... 1 jon fac 203 Feb 10 10:47 test
```

```
Bill> /usr/jon/test
```

has privilege of 'jon'

```
shell> chmod g+s
```

Has privilege of group 'fac'

Masks

```
creat ("my_file", 0777);
```

- Expectation:

- ```
shell> ls -l
```
- ```
-rwxrwxrwx jon ... my_file
```

- Instead:

- ```
-rw----- my_file
```

- What happened?

- To prevent against accidental exposure, Unix sets a default **mask** with your process (type `umask`)
- Typically: 077 (1 means mask out)

```
creat ("name", PERM & (~mask)); AND
regular files also mask out execute
```

## Masks (cont'd)

```
creat ("name", PERM & (~mask));
```

umask is 022: what is this one?

**To change the mask:**

```
mode_t umask (mode_t newmask);
```

# CSci 4061

## Introduction to Operating Systems

### **File Systems: Directories**

### **Chapter 5**


## Directory

- What is it?

# Directory

- Abstraction
  - Container for related files (and other directories)
- name
- location
- contents
- attributes

## Unix Path Names



A Unix directory tree  
MAX\_PATH: 1024 chars

## Path Names (cont'd)

- Home directory: dir you are logged into (~)
- Current working directory (cwd): `cd /usr/jim`

- `shell> pwd`

- `shell> /usr/jim`

- Relative file names(w/r to cwd)


- `shell> ls foo`

... advantage?

- Absolute file names (rooted from /)

- `Shell> ls /usr/jim/foo`

... advantage?



## Path Names (cont'd)

- Value of absolute path names:

`foo.c/foo`

...

```
f = fopen ("bar", "r");
```

...

Bill> `/usr/jon/foo` will fail unless "bar" is in cwd  
"bar" must be in the cwd of whomever runs it, instead:


```
f = fopen ("/user/jon/bar", "r");
```

Bill> `/usr/jon/foo` works now

On the other hand, if we were distributing `foo` ...

## Path Names (cont'd)

```
int chdir (const char *path);
```

  
**shell>** cd foo

Ex:

```
fd1 = open ("/usr/ben/abs", O_RDONLY);
chdir ("/usr/ben");
fd1 = open ("abs", O_RDONLY);
```

```
char *getcwd (char *name,
 size_t size);
```


  
**shell>** pwd

## Links: files with multiple names

Each name is an alias or a “hard link”

```
#include <unistd.h>
int link (const char *original_path,
 const char *new_path)
```

cannot exist as  
a file already



```
link ("foo", "bar"); // “bar” refers to file “foo”
unlink ("bar");
```

Number of links is the link count


Last unlink will delete the file (when no fd's to it)

Cannot unlink a directory

Does link affect the fd table?

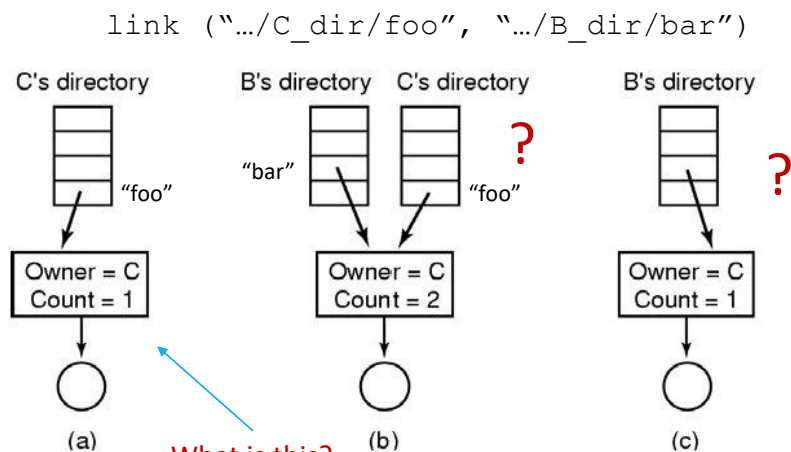


# Directories and Hard Links



File system containing a shared file

## Directories and Hard Links (cont'd)



```
unlink (".../C_dir/foo")
```

- a) Situation prior is linking
- b) After the link is created
- c) After the original owner unlinks/removes the file

# Symbolic Links

- Hard links cannot be made to directories or to files in other file systems
- Symbolic link: allows a file/dir name to “point to” another file/dir name

```
int symlink (const char *realname,
 const char *symname);
symlink ("/usr/jon/tmp1",
 "/usr/bill/tmp2");
```

New inode created for symname /usr/bill/tmp2

## Symbolic Links (cont'd)

```
symlink ("/user/jon/tmp1/f1", "f2");
lrwxrwxrwx f2 -> /usr/jon/tmp1/f1
```

Remove f2, symbolic link goes away, file does not

Remove /user/jon/tmp1/f1, symbolic link remains!

<board>

# Default Hard Links

- Two links: `.` and `..` (`ls -id <dir>`)
    - `.` Refers to `cwd`
    - `..` Refers to one level up from `cwd`
- ```
shell> ./cat
```

```
shell> cd /usr/jim/tmp
```

```
shell> ls ./foo
```

- same as `foo` or `/usr/jim/tmp/foo`

```
shell> ls ../bar
```

- same as `/usr/jim/bar`

Directory Permissions

- Directories are themselves represented by files
 - Have a name
 - Contents are file names
 - Same protection bits are used for directories (`rwX`)

Directory Permissions (cont'd)

- Read means class of users can list 'ls' contents of directory
- Write means class of users can create or remove files in the directory
- Execute means class of users can 'cd' into directory also allows open and execute for files in the directory

Directory Operations

- create/remove
- opendir/closedir
- readdir

Create

```
#include <sys/stat.h>
```

```
int mkdir (const char *pathname,  
           mode_t mode);
```

```
mkdir ("tmp/dir1", 0777);
```

Also places two links (. and .. in directory)

Remove

```
int rmdir (const char *pathname);
```

Removes the directory: directory must be empty!

Can be executed in the shell as well

```
shell> rmdir foo
```

Open/Close Directory

- Open a directory to look at its contents

```
#include <dirent.h>
```

```
DIR *opendir (const char *dirname);  
struct dirent *readdir (DIR *dirptr);  
int closedir (DIR *dirptr);
```

```
DIR *dp;  
dp = opendir ("/tmp/dir1");
```

```
struct dirent {  
    ino_t d_ino;  
    char d_name[NAMESIZE];  
}
```

readdir

```
#include <dirent.h>  
struct dirent *readdir  
                (DIR *dirptr);
```

Returns each directory entry, NULL at the end

readdir: example

Example: (very simple) my_ls

```
int my_ls (const char *name) {
    struct dirent *d;
    DIR *dp;
    if ((dp = opendir (name)) == NULL)
        return -1;
    while (d = readdir (dp))
        printf ("%s\n", d->d_name);
    closedir (dp);
    return 1;
}
```